



Member-only story

Kafka + Springboot Bootcamp — From Zero to Production



Arvind Kumar

Following

4 min read · Oct 29, 2025

110

4



...

Duration: 15 Days (1 hour/day)

Focus: Kafka Core + Spring Boot Integration + Real Production System

Target Audience: Developers new to Kafka but familiar with Java & Spring Boot

Apache Kafka isn't just a message broker — it's the **nervous system of modern distributed systems**.

From event-driven microservices to real-time analytics, Kafka sits at the center of systems that move data *continuously* instead of *occasionally*.

This **15-day Kafka Bootcamp** is designed to take a Java/Spring Boot developer from *zero* to *production-grade* Kafka integration — understanding the

internals, building fault-tolerant producers and consumers, handling retries, ensuring exactly-once delivery, and even understanding the Kafka deployed to Cloud.

[Full story for non-members](#) | [Grab My Microservices E-Book](#) | [Youtube](#) | [LinkedIn](#) | [Book a 1:1 Meeting](#)

The image is a promotional graphic for a Kafka Mastery bootcamp. It features a dark background with a central orange and white circular icon resembling a Kafka cluster. At the top left is a red rounded rectangle containing the text "PART 1". At the top right is a green rounded rectangle containing the text "Java + Spring Boot" next to a small icon. Below the icon is the text "BOOTCAMP SERIES". The main title "Kafka Mastery" is displayed in large, bold, orange and white letters. A horizontal bar below the title contains the text "From Zero to Production". Underneath the title, there are three rectangular boxes: "DURATION 15 Days", "DAILY TIME 1 Hour", and "LEVEL Beginner+". Below these boxes are three ovals: "HANDS-ON", "PRODUCTION-READY", and "REAL SYSTEMS".

Each day, we'll invest just **1 focused hour** — combining theory, guided demos, and hands-on mini assignments — to build up a solid Kafka-based system step by step.

Day 1: The World of Event Streaming

- Why Kafka? Evolution from message queues to event logs.

- Core concepts: topics, partitions, brokers, producers, consumers.
- Kafka architecture overview (ISR, leader-follower replication).

Hands-on: *Set up Kafka locally (using Docker or other way) and create your first topic.*

Day 2: CLI Practice and Deep Dive into Kafka Internals. Kafdrop for kafka monitoring

- Using `kafka-topics.sh`, `kafka-console-producer`, and `kafka-console-consumer`.
- Exploring topic metadata, partition leaders, and offsets.
- Understanding message durability and replication.
- Install Kafdrop for Kafka monitoring

Hands-on: *Produce and consume messages using CLI; simulate broker failure. Kafka monitoring with kafdrop*

Day 3: Spring Boot Integration — Producer Setup

- Setting up Kafka dependencies in Spring Boot.
- Producer configs: serializers, acks, retries, batching.
- Writing your first producer and publishing events programmatically.

Hands-on: *Send `orderPlacedEvent` using REST endpoint.*

Day 4: Spring Boot Integration — Consumer Setup

- Consumer groups and rebalancing explained.

- Deserialization and concurrent consumers.
- Manual vs auto offset commits.

| Hands-on: *Build an OrderConsumer that logs and persists data.*

Day 5: Error Handling + Retries + Dead Letter Topics

- Common failure patterns (serialization, transient errors, poison pills).
- Spring Kafka's `@RetryableTopic` and DLT setup.
- Logging and tracing failed messages.

| Hands-on: *Create a retry + DLT pipeline and verify with test messages.*

Day 6: Message Ordering and Keying Strategies

- Partitioning logic and message ordering guarantees.
- How keys influence message placement.
- Handling reordering in consumers.

| Hands-on: *Experiment with keyed vs keyless messages.*

Day 7: Idempotency, Transactions, and Exactly Once

- Producer idempotence (`enable.idempotence=true`).
- Kafka transactions and atomic writes to multiple topics.
- Exactly-once semantics — myth vs reality.

| Hands-on: *Build a payment processor with transactional producer.*

Day 8: Building Resilient Event-Driven Systems

- Event delivery patterns: fire-and-forget, synchronous, and async acknowledgment.
- Designing event contracts and versioning strategies.
- Real-world pitfalls: duplicate messages, network retries, and lag.

| Hands-on: *Add logging and validation layers around events.*

Day 9: Outbox Pattern

- Problem: dual writes and data consistency.
- Outbox table design, transactional publishing, and Debezium overview.
- Implementing the outbox pattern with Spring Boot + Kafka.

| Hands-on: *Implement outbox publisher for Order Service.*

Day 10: Saga Pattern and Event Choreography

- Coordinating distributed transactions with Kafka events.
- Saga orchestrator vs choreography models.
- Example: Order → Payment → Inventory → Notification.

| Hands-on: *Implement event-based Saga flow.*

Day 11: Kafka Streams Fundamentals

- Stream processing vs message consumption.
- Stateless transformations (`map`, `filter`, `branch`).

- Writing a Kafka Streams app inside Spring Boot.

| Hands-on: *Build order event transformer service.*

Day 12: Kafka Streams — Stateful Operations

- Aggregations, joins, and windowed computations.
- Understanding state stores and RocksDB.
- Reprocessing and materialized views.

| Hands-on: *Count orders per minute and expose via REST.*

Day 13: Schema Registry + Avro + Versioning

- Why schemas matter in event-driven systems.
- Using Confluent Schema Registry.
- Avro serialization and compatibility rules (backward, forward).

| Hands-on: *Produce and consume Avro messages.*

Day 14: Integration Testing with Embedded Kafka

- Unit testing producers and consumers.
- Integration testing with `@EmbeddedKafka`.
- Testcontainers for Kafka end-to-end flow.

| Hands-on: *Write integration tests for event pipeline.*

Day 15: Monitoring, Lag Metrics, and Cloud Deployment

- Understanding consumer lag and offsets.
- Kafka metrics via Micrometer and Actuator.
- Understanding Kafka available on AWS (known as MSK) and Confluent Cloud.
- Production tuning tips (acks, linger.ms, batch.size).

| Hands-on: *Deploy and verify your system's metrics dashboard.*

Final Project

Goal: Build a mini “Order Management System” consisting of:

- **Order Service** → produces events
- **Payment Service** → consumes + produces
- **Inventory Service** → consumes + updates stock
- Add retry/DLT handling, transaction safety, and test coverage.

Closing Note

By the end of this 15-day bootcamp, you won't just “know Kafka” — You will understand *why* Kafka behaves the way it does under load, during failures, and across distributed systems.

They'll be able to:

- Build reliable producers and consumers in Spring Boot.

- Handle retries, DLTs, and transactional events with confidence.
- Test and monitor Kafka pipelines like a production engineer.
- Deploy and operate Kafka apps on cloud environments such as MSK or Confluent Cloud.

| *Need my help to do this bootcamp? Drop message on [linkedin](#)*

This lays the perfect foundation for **Part 2: Advanced Kafka Master class**, where we'll dive into **Kafka Connect, Schema Evolution, Performance Tuning, Multi-Cluster Replication, and Real-Time Analytics Pipelines** — the topics that turn an engineer into an architect.

=====

Here is the list for all the stories around Kafka Boot Camp

The screenshot shows a Medium article card. At the top left is a profile picture of Arvind Kumar. To the right of the profile picture is the author's name, Arvind Kumar. Below the author information is the title of the article, "Kafka Bootcamp | Master Kafka with Codefarm". Underneath the title is a button labeled "View list". To the right of the button is the number "15 stories". At the bottom of the card is a small footer with the text "APACHE KAFKA • SPRING BOOT".

11: Kafka Stream Fundamentals

APACHE KAFKA • SPRING BOOT

Kafka

Apache Kafka

Kafka Producer

Kafka Streams

Kafka Connect



Written by Arvind Kumar

3.2K followers · 110 following

Following ▾



Staff Engineer | System Design, Microservices, Java, SpringBoot, Kafka, DBs, AWS, GenAI | Teaching concepts via stories & characters | linkedin.com/in/codefarm0

Responses (4)



Pradeep Kasula

What are your thoughts?



Enamoradoc

Oct 29, 2025 (edited)

...

Awesome, glad you're doing this!



7



1 reply

[Reply](#)



Engouyassa

Nov 11, 2025

...

Great undertaking. How is this bootcamp accessible?



4



1 reply

[Reply](#)



mohamad shahkhajeh

Nov 2, 2025

...

Have you considered how partitioning strategy might impact exactly-once delivery and consumer lag under load?

[Reply](#)[See all responses](#)

More from Arvind Kumar



In FAUN.dev() by Arvind Kumar

The Silent Earthquake in Tech: Why Every Engineer Must Rethink Thei...

Sharing my personal experience, thought process, and happenings around me.



Dec 13, 2025



482



20



Arvind Kumar

Building a Real-Time Messaging System: Design Deep Dive

How do you deliver messages to 2 billion users in real-time? Let's design a messaging...



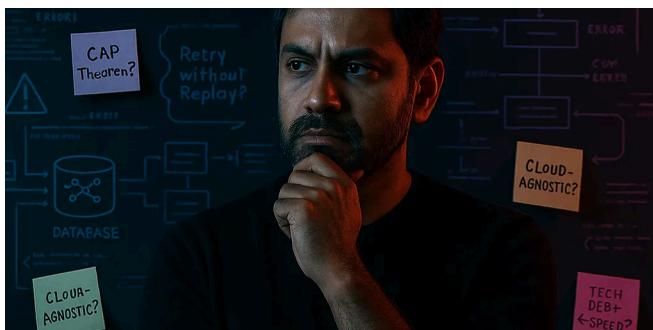
Dec 7, 2025



131



2





Arvind Kumar

12 Interview Questions That Separate Real Architects from...

“Staff Engineer or Just Staff Meeting Engineer?”



Jul 3, 2025



121



1



Arvind Kumar

Forward Proxy vs Reverse Proxy: The Deep Dive Every Engineer Mu...

There are few networking components as critical—and as misunderstood—as forward...



Dec 4, 2025



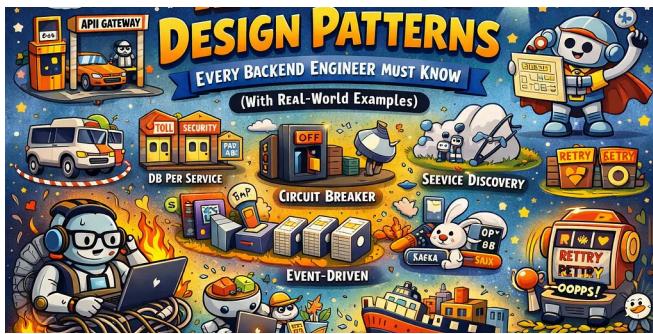
188



2

[See all from Arvind Kumar](#)

Recommended from Medium



KAFKA + SPRINGBOOT BOOTCAMP: LEARN BY BUILDING

Day 9: Outbox Pattern

APACHE KAFKA • SPRING BOOT

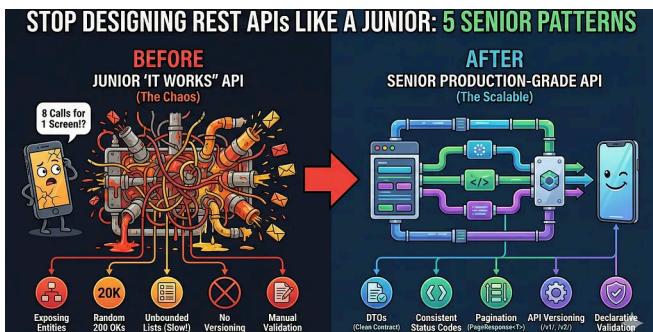
Gaddam.Naveen

If You Know These 12 Microservice Design Patterns, You're Interview-...

if you are not a medium member then Click here to read free

Dec 20, 2025 14 2

4d ago 3



In Stackademic by Habibwahid

Stop Designing REST APIs Like a Junior: 5 Patterns Senior Engine...

Your API works. But is it maintainable, scalable, and client-friendly? Let's understand...

Dec 10, 2025 178 1

Code With Sunil | Code Smarter, not harder

Spring Boot Errors, Exceptions, and AOP Annotations—A Practical...

Learn how to centralize error handling, simplify exception management, and use AO...

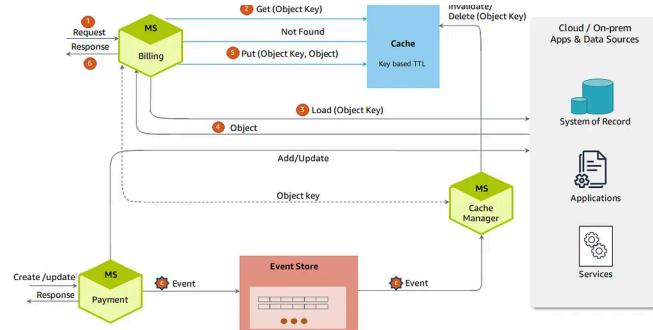
4d ago 51

```

34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
      self.debug = debug
      self.logger = logging.getLogger(__name__)
      if path:
          self._file = open(os.path.join(path))
          self._file.seek(0)
          self._fingerprints = set()
      else:
          self._file = None
      self._fp_size = 1024
      self._fp_offset = 0
      self._fp_index = 0
      self._fp_line = b''

      @classmethod
      def from_settings(cls, settings):
          debug = settings.getboolean('superuser_debug')
          return cls(job_dir=settings), debug

      def request_seen(self, request):
          fp = self.request_fingerprint(request)
          if fp in self._fingerprints:
              return True
          self._fingerprints.add(fp)
          if self._file:
              self._file.write(fp + os.linesep)
  
```



In CodeX by Riya Sharma

2 Ways to Make Spring Boot Fully Serverless in 2026

A practical guide to running Spring Boot without servers, ops headaches, or idle costs

5d ago 83

In DevOps.dev by StackStories

Caching Strategy for Microservices in Spring Boot (Production-Ready...)

Caching in microservices is not optional at scale. Without a well-designed caching...

Dec 26, 2025 13 1

[See more recommendations](#)