

Member-only story

The World of Event Streaming — Day 1

Welcome to Kafka!



Arvind Kumar

Following

9 min read · Dec 22, 2025

70

4



...

Imagine you're at a busy restaurant. The old way (message queues) is like a waiter taking your order, delivering it to the kitchen, waiting for the food, and bringing it back. One order at a time, slow and sequential.

Kafka is like a **conveyor belt system** in a modern restaurant. Orders come in continuously, get placed on the belt, and multiple chefs can grab orders, customers can see what's available, and managers can track everything in real-time. Everyone works at their own pace, and nothing gets lost.

That's event streaming — data flows continuously, and multiple systems can consume it independently.

[Full story for non-members](#) | [Microservices E-Book](#) | [Spring Boot E-Book](#) | [Join Whatsapp Group for Daily Tech Bytes](#) | [Youtube](#) | [LinkedIn](#)



Why Kafka? Evolution from Message Queues to Event Logs

The Old Way: Message Queues

Traditional message queues (like RabbitMQ, ActiveMQ) work like a **postal system**:

- You send a letter (message) to a recipient
- Once delivered and acknowledged, the letter is typically removed
- Messages are “consumed” — after reading, they’re gone

- RabbitMQ CAN send to multiple consumers (via fanout/topic exchanges), but messages are usually deleted after acknowledgment

Key Characteristics:

- Messages are delivered and removed (or moved to dead-letter queue)
- No built-in replay capability – once consumed, you can't re-read it
- Designed for task distribution and request/response patterns
- Great for: RPC, task queues, request routing

The New Way: Event Logs (Kafka)

Kafka works like a **newspaper archive**:

- Events are published (like articles) and stored permanently
- Multiple readers can subscribe and read independently
- You can read today's paper, yesterday's, or even last week's
- New readers can join anytime and catch up on history
- Events are retained for a configurable period (days, weeks, or forever)

Key Differences:

- **Message Queue:** “Deliver this message” → consumed → deleted (postal system)
- **Event Log:** “Publish this event” → stored → replayable by anyone (newspaper archive)

Why This Matters:

- **Replay capability:** Re-process events if your consumer crashes or needs to catch up
- **Multiple independent consumers:** Each consumer reads at their own pace from their own position
- **Event sourcing:** Build state by replaying all events from the beginning
- **Audit trail:** All events are preserved for compliance and debugging

Real-World Example

Think of an e-commerce system:

- **Old way:** Order service sends message to Payment service → waits → sends to Inventory service → waits
- **Kafka way:** Order service publishes “OrderPlaced” event → Payment, Inventory, Email, Analytics all consume it simultaneously, at their own pace

Core Concepts: Building Blocks of Kafka

Let's understand Kafka's core concepts with simple analogies:

1. Topic

A topic is like a **newspaper section** (Sports, Business, Technology).

- All related events go into one topic

- Example: orders, payments, user-events

2. Partition

A partition is like a **page** in that newspaper section.

- Topics are split into partitions for parallel processing
- Each partition is an ordered sequence of events
- Think of it as multiple lanes on a highway — more lanes = more throughput

3. Broker

A broker is like a **newspaper printing press**.

- Kafka cluster = multiple printing presses (brokers)
- Each broker stores some partitions
- If one press breaks, others keep printing

4. Producer

A producer is like a **reporter** writing articles.

- Applications that publish events to topics
- Example: Order Service producing “OrderPlaced” events

5. Consumer

A consumer is like a **reader** subscribing to the newspaper.

- Applications that read events from topics
- Example: Payment Service consuming “OrderPlaced” events

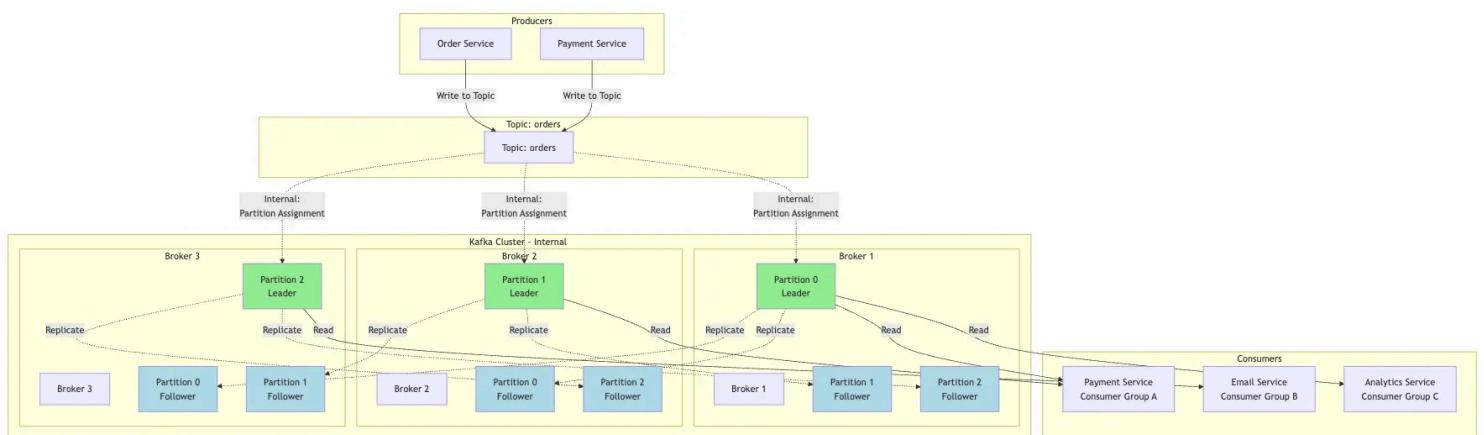
6. Consumer Group

A consumer group is like a family sharing one newspaper subscription.

- Multiple consumers in a group share the work
- Each partition is read by only one consumer in the group
- If one consumer is slow, others pick up the slack

Kafka Architecture Overview

Visual Architecture Diagram



Legend:

- **Green (Leader):** Handles all read/write requests

- **Light Blue (Follower):** Replicates data from leader
- **Solid arrows:** Producer/Consumer operations (external)
- **Dashed arrows:** Internal Kafka operations (partition assignment & replication)

Leader-Follower Replication

Imagine a master chef and assistant chefs:

- **Leader:** The master chef (handles all requests)
- **Followers:** Assistant chefs (copy everything the master does)
- If the master chef gets sick, an assistant chef becomes the new master

In Kafka:

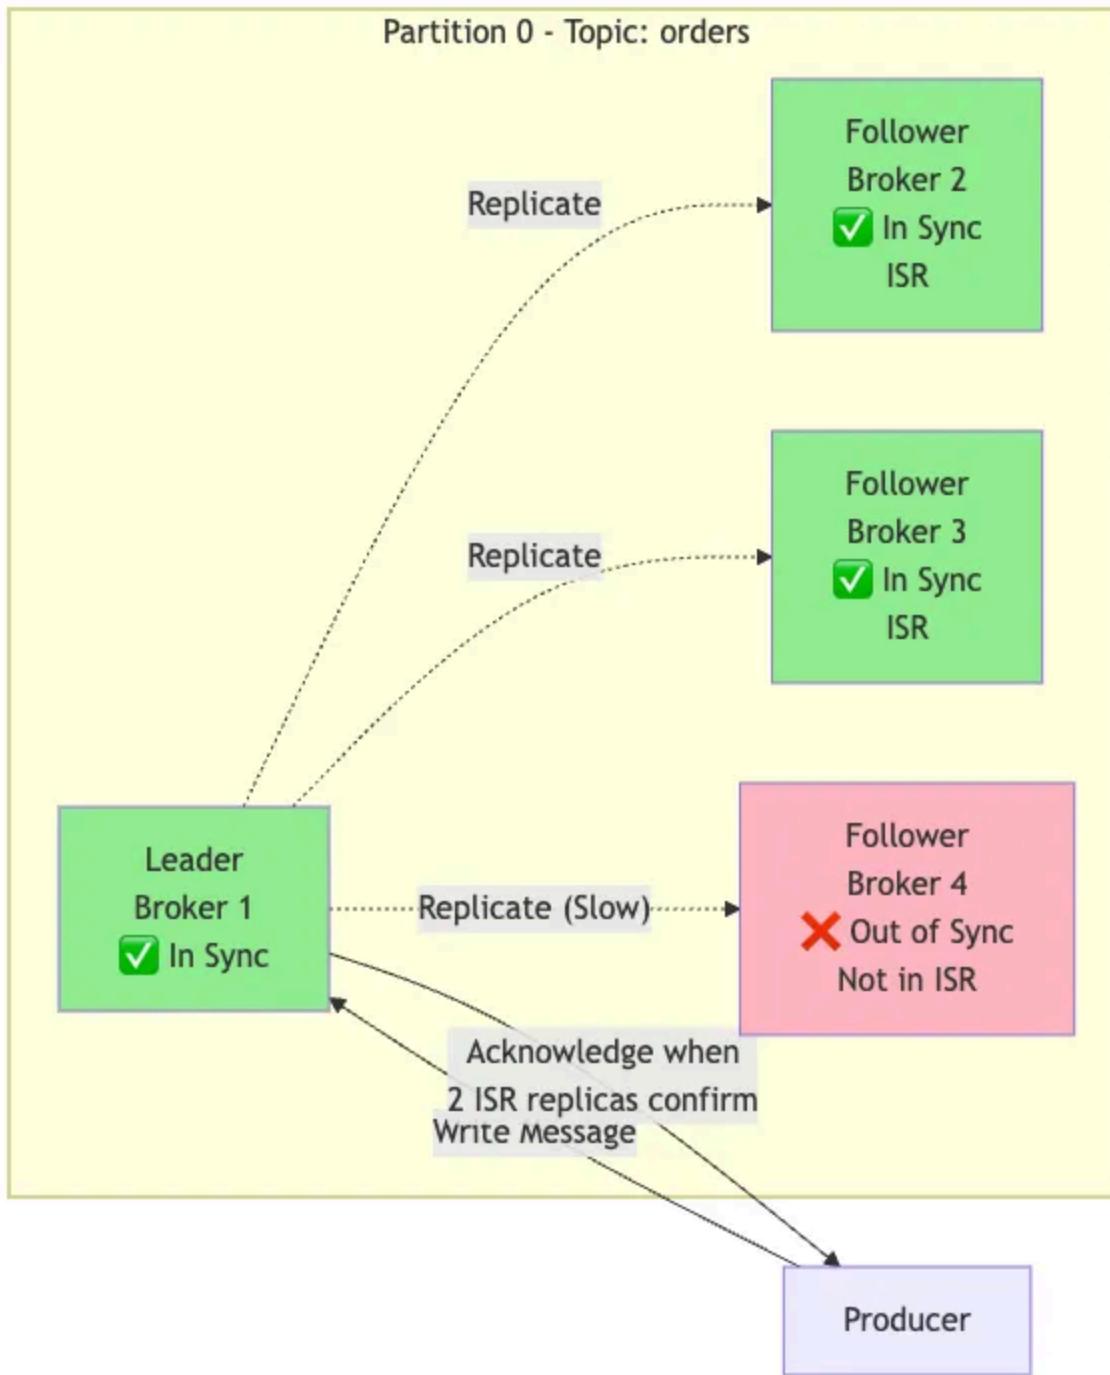
- Each partition has one **leader** (handles reads/writes)
- Multiple **followers** (replicate data)
- If leader fails, a follower becomes the new leader automatically

ISR (In-Sync Replicas)

ISR is like a quality control team:

- Only followers that are “caught up” with the leader are in ISR
- If a follower falls behind (network issue, slow disk), it’s removed from ISR
- Leader only considers writes successful when ISR replicas confirm

Why this matters: Ensures data durability. Even if a broker crashes, your data is safe on other brokers.



What this shows:

- **Leader and 2 Followers are in ISR (In-Sync Replicas)** — they're caught up

- 1 Follower is out of sync (maybe network issue or slow disk) — removed from ISR
- Producer gets acknowledgment when **ISR replicas** confirm (not the out-of-sync one)
- If leader crashes, only ISR followers can become the new leader

KRaft Mode (No More Zookeeper!)

Modern Kafka (4.1.0+) uses KRaft mode:

- **Old way:** Kafka relied on Zookeeper for cluster coordination and metadata storage
- **New way:** Kafka uses its own **KRaft** (Kafka Raft) protocol for cluster management
- **Benefits:** Simpler setup, better performance, easier scaling, no separate Zookeeper cluster needed

Think of it like a company that used to hire an external manager (Zookeeper) but now manages itself internally (KRaft). Everything runs smoother and faster!

Hands-on: Setting Up Kafka Locally

Let's get Kafka running on your machine. We'll use Docker (the easiest way) with official Apache Kafka 4.1.0 (open source).

Good news! Modern Kafka (4.1.0+) no longer needs Zookeeper! It uses KRaft mode (Kafka Raft) for cluster management, making setup simpler.

Prerequisites

- Docker Desktop installed and running
- Basic terminal/command line knowledge

If you want to run kafka without Docker follow this guide — [link](#)

Step 1: Create Docker Compose File

Create a file named `docker-compose.yml` in a new folder. We're using the latest official Apache Kafka with KRaft mode (no Zookeeper needed):

```
services:  
  kafka:  
    image: apache/kafka:4.1.0  
    container_name: kafka  
    ports:  
      - "9092:9092"  
    environment:  
      KAFKA_PROCESS_ROLES: broker,controller  
      KAFKA_NODE_ID: 1  
      KAFKA_CONTROLLER_QUORUM_VOTERS: 1@kafka:9093  
  
      KAFKA_LISTENERS: PLAINTEXT://0.0.0.0:9092,PLAINTEXT_INTERNAL://0.0.0.0:290  
      KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://localhost:9092,PLAINTEXT_INTERNAL:  
      KAFKA_LISTENER_SECURITY_PROTOCOL_MAP: CONTROLLER:PLAINTEXT,PLAINTEXT:PLAIN  
      KAFKA_CONTROLLER_LISTENER_NAMES: CONTROLLER  
  
      KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1  
      KAFKA_TRANSACTION_STATE_LOG_REPLICATION_FACTOR: 1  
      KAFKA_TRANSACTION_STATE_LOG_MIN_ISR: 1  
  
      KAFKA_LOG_DIRS: /var/lib/kafka/data  
      CLUSTER_ID: MkU30EVBNTcwNTJENDM2Qk  
  
    volumes:
```

```
- kafka_data:/var/lib/kafka/data

kafdrop:
  image: obsidiandynamics/kafdrop:latest
  container_name: kafdrop
  ports:
    - "9000:9000"
  environment:
    KAFKA_BROKERCONNECT: kafka:29092
  depends_on:
    - kafka

volumes:
  kafka_data:
```

What's different?

- **No Zookeeper!** Kafka manages itself using KRaft
- **KAFKA_PROCESS_ROLES: broker,controller** - Single node acts as both broker and controller
- **KAFKA_CONTROLLER_QUORUM_VOTERS** - Defines the controller quorum (just one node for local dev)

Step 2: Start Kafka

Open your terminal in the folder where you created `docker-compose.yml` and run:

```
docker-compose up -d
```

```
codefarm@Arvinds-MacBook-Pro notification-service % docker-compose up -d
[+] Running 3/3
  ✓ Network notification-service_default  Created
  ✓ Container kafka                      Started
  ✓ Container kafdrop                   Started

codefarm@Arvinds-MacBook-Pro notification-service %
```

This starts Kafka in KRaft mode. Wait about 30–45 seconds for it to fully start (Kafka needs to format the storage on first run).

Verify it's running:

```
docker-compose ps
```

```
codefarm@Arvinds-MacBook-Pro notification-service % docker-compose ps
NAME      IMAGE               COMMAND             SERVICE   CREATED        STATUS     PORTS
kafdrop   obsidiandynamics/kafdrop:latest "/kafdrop.sh"    kafdrop   15 seconds ago Up 15 seconds  0.0.0.0:9000->9000/tcp, [::]:9000->9000/tcp
kafka     apache/kafka:4.1.0   "/__cacert_entrypoin..." kafka    15 seconds ago Up 15 seconds  0.0.0.0:9092->9092/tcp, [::]:9092->9092/tcp
```

You should see the `kafka` service running. Check the logs to ensure it started successfully:

```
docker-compose logs kafka
```

```
kafka | [2025-12-28 10:19:26,791] INFO authorizerStart completed for endpoint PLAINTEXT. Endpoint is now READY. (org.apache.kafka.common.protocol.Authorizer)
kafka | [2025-12-28 10:19:26,791] INFO authorizerStart completed for endpoint PLAINTEXT_INTERNAL. Endpoint is now RE/started (org.apache.kafka.common.protocol.Authorizer)
kafka | [2025-12-28 10:19:26,791] INFO [SocketServer listenerType=BROKER, nodeId=1] Enabling request processing. (kafka.network.SocketServer)
kafka | [2025-12-28 10:19:26,791] INFO Awaiting socket connections on 0.0.0.0:9092. (kafka.network.DataPlaneAcceptor)
kafka | [2025-12-28 10:19:26,792] INFO Awaiting socket connections on 0.0.0.0:29092. (kafka.network.DataPlaneAcceptor)
kafka | [2025-12-28 10:19:26,795] INFO [BrokerServer id=1] Waiting for all of the authorizer futures to be completed (kafka.server.BrokerServer)
kafka | [2025-12-28 10:19:26,795] INFO [BrokerServer id=1] Finished waiting for all of the authorizer futures to be completed (kafka.server.BrokerServer)
kafka | [2025-12-28 10:19:26,795] INFO [BrokerServer id=1] Waiting for all of the SocketServer Acceptors to be started (kafka.server.BrokerServer)
kafka | [2025-12-28 10:19:26,795] INFO [BrokerServer id=1] Finished waiting for all of the SocketServer Acceptors to be started (kafka.server.BrokerServer)
kafka | [2025-12-28 10:19:26,795] INFO [BrokerServer id=1] Transition from STARTING to STARTED (kafka.server.BrokerServer)
kafka | [2025-12-28 10:19:26,795] INFO Kafka version: 4.1.0 (org.apache.kafka.common.utils.AppInfoParser)
kafka | [2025-12-28 10:19:26,797] INFO Kafka commitId: 13f70256db3c994c (org.apache.kafka.common.utils.AppInfoParser)
kafka | [2025-12-28 10:19:26,797] INFO Kafka startTimeMs: 1766917166795 (org.apache.kafka.common.utils.AppInfoParser)
kafka | [2025-12-28 10:19:26,797] INFO [KafkaRaftServer nodeId=1] Kafka Server started (kafka.server.KafkaRaftServer)
kafka | [2025-12-28 10:19:26,880] INFO [GroupCoordinator id=1] Scheduling loading of metadata from __consumer_offsets-13 with epoch 4 (org.apache.kafka.coordinator.GroupCoordinator)
kafka | [2025-12-28 10:19:26,880] INFO [GroupCoordinator id=1] Scheduling loading of metadata from __consumer_offsets-46 with epoch 4 (org.apache.kafka.coordinator.GroupCoordinator)
```

Look for messages like “Kafka Server started” — this means Kafka is ready!

Step 3: Access Kafdrop (Kafka Web UI)

Kafdrop is a web-based UI for monitoring your Kafka cluster. It's already included in the docker-compose file above!

Access Kafdrop:

1. Open your web browser
2. Navigate to: <http://localhost:9000>
3. You should see the Kafdrop interface

What you'll see:

- **Topics list:** All topics in your cluster (empty for now)
- **Brokers:** Information about your Kafka brokers
- **Consumer groups:** Active consumer groups
- **Message browser:** View messages in topics (once you create some)

The screenshot shows the Kafdrop Kafka Cluster Overview page. At the top, there's a navigation bar with icons for back, forward, search, and other browser functions, and the URL 'localhost:9000'. The title 'Kafdrop' is displayed with its logo, and a 'Star' button. Below the title, it says '4.2.0 [2025-07-31T09:48:58.619Z]'. The main content area is titled 'Kafka Cluster Overview'.

Bootstrap servers:

Bootstrap servers	kafka:29092
-------------------	-------------

Total topics:

Total topics	0
--------------	---

Total partitions:

Total partitions	0
------------------	---

Total preferred partition leader:

Total preferred partition leader	0%
----------------------------------	----

Total under-replicated partitions:

Total under-replicated partitions	0
-----------------------------------	---

Brokers:

ID	Host	Port	Rack	Controller	Number of partitions (% of total)
1	kafka	29092	-	Yes	0 (0%)

Topics:

Name	Partitions	% Preferred	# Under-replicated	Custom Config
No topics available	(0)			

ACLs:

+ New

Kafdrop features:

- ✓ Browse topics and partitions
- ✓ View messages in real-time
- ✓ See consumer group lag
- ✓ Monitor broker health
- ✓ Create topics (though we'll use CLI for learning)

Note: If Kafdrop doesn't load immediately, wait 10–20 seconds for it to start and connect to Kafka.

Hands-on: Create Your First Topic

Now let's create a topic! Think of this as creating a new newspaper section.

Step 1: Enter the Kafka Container

```
docker exec -it <kafka-container-name> bash
```

```
codefarm@Arvinds-MacBook-Pro notification-service % docker exec -it kafka bash  
cc57d0c5d04b:/$
```

To find the container name:

```
docker ps
```

Look for the container with `apache/kafka` image. Copy its name (it should be `kafka` if you used the docker-compose above, or something like.

Step 2: Create a Topic

Inside the container, run:

Move to bin directory `:/opt/kafka/bin` and check all commands available

```
cc57d0c5d04b:/opt/kafka/bin$ ls
connect-distributed.sh      kafka-console-consumer.sh      kafka-features.sh      kafka-reassign-partitions.sh      kafka-streams-groups.sh
connect-mirror-maker.sh     kafka-console-producer.sh    kafka-get-offsets.sh    kafka-replica-verification.sh    kafka-topics.sh
connect-plugin-path.sh      kafka-console-share-consumer.sh  kafka-groups.sh        kafka-run-class.sh          kafka-transactions.sh
connect-standalone.sh       kafka-consumer-groups.sh   kafka-jmx.sh           kafka-server-start.sh       kafka-verifiable-consumer.sh
kafka-acls.sh              kafka-consumer-perf-test.sh  kafka-leader-election.sh  kafka-server-stop.sh         kafka-verifiable-producer.sh
kafka-broker-api-versions.sh kafka-delegation-tokens.sh  kafka-log-dir.sh       kafka-share-consumer-perf-test.sh  kafka-verifiable-share-consumer.sh
kafka-client-metrics.sh    kafka-delete-records.sh    kafka-metadata-quorum.sh  kafka-share-groups.sh        trogrodor.sh
kafka-cluster.sh           kafka-dump-log.sh        kafka-metadata-shell.sh  kafka-storage.sh          windows
kafka-configs.sh           kafka-e2e-latency.sh    kafka-producer-perf-test.sh  kafka-streams-application-reset.sh
```

Then command to create the topic

```
./kafka-topics.sh --create \
--bootstrap-server localhost:9092 \
--replication-factor 1 \
--partitions 3 \
--topic my-first-topic
```

```
cc57d0c5d04b:/opt/kafka/bin$ ./kafka-topics.sh --create \
--bootstrap-server localhost:9092 \
--replication-factor 1 \
--partitions 3 \
--topic my-first-topic
Created topic my-first-topic.
cc57d0c5d04b:/opt/kafka/bin$ █
```

What this means:

- `--create` : We're creating a new topic
- `--bootstrap-server localhost:9092` : Where Kafka is running
- `--replication-factor 1` : How many copies of data (1 = no replication, fine for local dev)
- `--partitions 3` : Split the topic into 3 partitions (like 3 lanes on a highway)

- --topic my-first-topic : The name of our topic

You should see: Created topic my-first-topic.

Step 3: Verify Your Topic

Let's list all topics to confirm it was created:

```
./kafka-topics.sh --list --bootstrap-server localhost:9092
```

You should see my-first-topic in the list.

```
cc57d0c5d04b:/opt/kafka/bin$ ./kafka-topics.sh --list --bootstrap-server localhost:9092  
my-first-topic  
cc57d0c5d04b:/opt/kafka/bin$ █
```

Step 4: Describe Your Topic (Optional but Useful)

See detailed information about your topic:

```
./kafka-topics.sh --describe \  
--bootstrap-server localhost:9092 \  
--topic my-first-topic
```

This shows:

- Partition count (3)

- Replication factor (1)
- Which broker is the leader for each partition

```
cc57d0c5d04b:/opt/kafka/bin$ ./kafka-topics.sh --describe \
--bootstrap-server localhost:9092 \
--topic my-first-topic
Topic: my-first-topic    TopicId: Lz3aK68jTfSt4jv2bV0EOA PartitionCount: 3      ReplicationFactor: 1   Configs: min.insync.replicas=1
    Topic: my-first-topic    Partition: 0    Leader: 1      Replicas: 1      Isr: 1   Elr:   LastKnownElr:
    Topic: my-first-topic    Partition: 1    Leader: 1      Replicas: 1      Isr: 1   Elr:   LastKnownElr:
    Topic: my-first-topic    Partition: 2    Leader: 1      Replicas: 1      Isr: 1   Elr:   LastKnownElr:
cc57d0c5d04b:/opt/kafka/bin$
```

Exit the container:

```
exit
```

```
cc57d0c5d04b:/opt/kafka/bin$ exit
exit

codefarm@Arvinds-MacBook-Pro notification-service %
```

What You've Learned Today

- ✓ Why Kafka? — Event logs vs message queues (newspaper archive vs postal system)
- ✓ Core Concepts — Topics, partitions, brokers, producers, consumers
- ✓ Architecture — Leader-follower replication and ISR
- ✓ Hands-on — Set up Kafka locally and created your first topic

Quick Check: Do You Understand?

Before moving to Day 2, make sure you can answer:

1. What's the difference between a message queue and an event log?
2. Why do we split topics into partitions?
3. What happens if a Kafka broker (leader) crashes?
4. What does ISR stand for and why is it important?

What's Next?

Tomorrow (Day 2), you'll:

- Use CLI tools to produce and consume messages
- Explore topic metadata and offsets
- Simulate broker failures
- Set up Kafdrop for visual monitoring

Congratulations! You've completed Day 1. You now have Kafka running locally and understand the fundamentals. Take a break, and we'll dive deeper tomorrow!

Troubleshooting

Kafka won't start?

- Make sure Docker Desktop is running
- Check if ports 9092 (broker) or 9093 (controller) are already in use
- Try: `docker-compose down` then `docker-compose up -d`
- On first run, Kafka needs to format storage — wait 30–45 seconds

Can't find the container?

- Run `docker ps -a` to see all containers (including stopped ones)
- Container name should be `kafka` (or `<folder-name>-kafka-1` if using default naming)

Topic creation fails?

- Wait 30–60 seconds after starting Kafka (longer on first run)
- Check logs: `docker-compose logs kafka`
- Look for “Kafka Server started” in the logs to confirm it’s ready

KRaft mode issues?

- If you see errors about CLUSTER_ID, make sure the environment variable is set correctly
- For a fresh start, remove the container and volumes: `docker-compose down -v`

=====

Below is a collection of all the stories in one place

List: Kafka Bootcamp | Master Kafka with Codefarm | Curated by Arvind Kumar | Medium

Kafka Bootcamp | Master Kafka with Codefarm · Kafka Bootcamp related stories · 1 stories on Medium

codefarm0.medium.com

Kafka

Event Driven Architecture

Apache Kafka

Kafka Consumer

Kafka Producer



Written by Arvind Kumar

3.2K followers · 110 following

Following ▾



Staff Engineer | System Design, Microservices, Java, SpringBoot, Kafka, DBs, AWS, GenAI | Teaching concepts via stories & characters | linkedin.com/in/codefarm0

Responses (4)



Pradeep Kasula

What are your thoughts?



Priyanshi
Dec 28, 2025

...

very well explained.



1



1 reply

[Reply](#)



San Rawat
Dec 27, 2025

...

why localhost is not working rather replace with kafka in yml file make it working.



1



1 reply

[Reply](#)



Kritikakhanna she/her
Dec 24, 2025

...

I loved how you managed to create a simple analogy like newspaper to describe components of Kafka. It is easy to follow and remember.



1



1 reply

[Reply](#)

[See all responses](#)

More from Arvind Kumar



In FAUN.dev() 🐾 by Arvind Kumar

The Silent Earthquake in Tech: Why Every Engineer Must Rethink Thei...

Sharing my personal experience, thought process, and happenings around me.

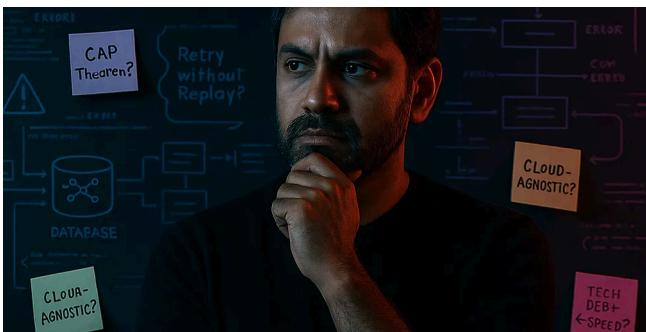
Dec 13, 2025 482 20

Arvind Kumar

Building a Real-Time Messaging System: Design Deep Dive

How do you deliver messages to 2 billion users in real-time? Let's design a messaging...

Dec 7, 2025 131 2



Arvind Kumar

12 Interview Questions That Separate Real Architects from...

"Staff Engineer or Just Staff Meeting Engineer?"

Jul 3, 2025 121 1



Arvind Kumar

Forward Proxy vs Reverse Proxy: The Deep Dive Every Engineer Mu...

There are few networking components as critical—and as misunderstood—as forward...

Dec 4, 2025 188 2

See all from Arvind Kumar

Recommended from Medium



Suchirreddy

Designing a Concurrency-Safe Ticket Booking System from...

Defining the Problem

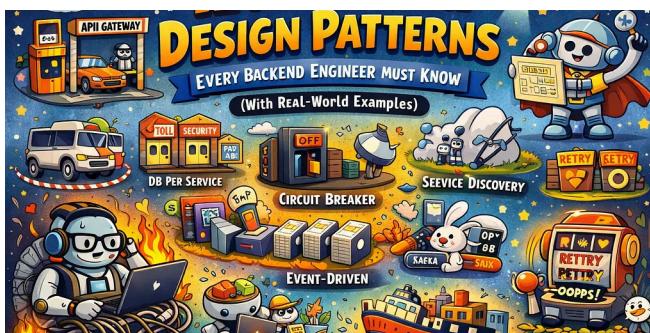
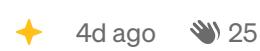
Oct 20, 2025



Arvind Kumar

Common Tricky Questions in the Kafka Ecosystem

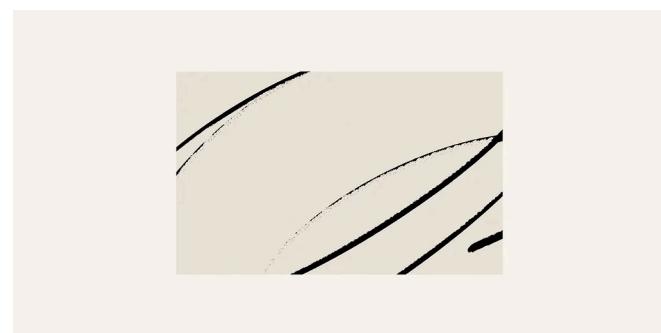
These common questions are part of the Kafka Bootcamp learning series—link



Gaddam.Naveen

If You Know These 12 Microservice Design Patterns, You're Interview...

if you are not a medium member then Click here to read free



Ram

Part 2: Master Kubernetes Interviews in 2026: Top 10 Essential...

Kubernetes has evolved far beyond basic container orchestration and, by 2026, it has...

Dec 20, 2025

14

2



...

6d ago

9



...



```

34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53

self.debug = debug
self.logger = logging.getLogger(__name__)
if path:
    self.file = open(os.path.join(path, 'logger.log'), 'a')
    self.file.seek(0)
    self.fingerprints = set()
    self._write_header()

@classmethod
def from_settings(cls, settings):
    debug = settings.getbool('DEBUGGING', debug)
    return cls(job_dir(settings), debug)

def request_seen(self, request):
    fp = self.request_fingerprint(request)
    if fp in self.fingerprints:
        return True
    self.fingerprints.add(fp)
    if self.file:
        self._write(fp + os.linesep)

def _write_header(self):
    if self.file:
        self._write(header)

```

CoVaib DeepLearn

Day 10 : System Design Concept : Event Driven Architecture

What Is Event-Driven Architecture?

Jul 26, 2025



...

5d ago

83



...

In CodeX by Riya Sharma

2 Ways to Make Spring Boot Fully Serverless in 2026

A practical guide to running Spring Boot without servers, ops headaches, or idle costs

[See more recommendations](#)